

myThread: User Transparent Multi-Processes Programming Library

Tianyi Liu

Computer Science Department
University of Texas at San Antonio
San Antonio, TX, USA
tianyi.liu@my.utsa.edu

Abstract

Multithreaded programming is an efficient way to enhance performance of applications. However, multithreaded programs always have data race problem, which makes the result of parallel programs be undeterministic. Undeterministic feature can be a fatal weakness for precise critical applications and even induce disasters in some cases. To deal with this problem, this technique report present `myThread`, a run time library, which could convert threads into processes silently, and the whole process is transparent to programmers.

When we use processes to implement threads, we should guarantee the processes can also share the memory regions that threads share in multithreaded programs. Also, we should reinitialize the mutex locks, condition variables and barriers to make them shared among processes. The experiment result indicates that `myThread` can be comparable with `pthread` and in some cases can be even better.

Keywords Multithreaded Programs, Multi-Processes, `pthread`, `Mmap`, Transparent

1 Introduction

Multithreading is an effective way to enhance performance of applications. It tries to fully use of the resources in modern computer system. However, Multithreaded programs also have to face some problems, such as data race, undeterministic executing flow, and deadlock. There are two reasons for such phenomena. The first reason is that multithreaded programs share global data and heap, and different threads can access others' stacks if they have the valid addresses. The second reason is that there is no guarantee of the executing order among different threads, and they may access the same memory locations with different orders.

Because of the undeterministic characteristic, the output of multi-threaded programs may vary if we run a program for several times, which is really a serious problem in precise sensitive computing tasks and finance or economy related affairs, and it may even induce disasters because of the undeterministic results.

There exist some systems to deal with the undeterministic problem. `CoreDet` [1] is a system to check every undeterministic operations, such as memory accesses, synchronizations. If the memory location or synchronization is owned by a thread itself, then the thread executes the command directly. Otherwise, it will wait and check again in next stage. `CoreDet` is really a useful design for avoiding undeterministic behaviors, but it is slow and unstable. `Dthread` [3] is another system designed to solve undeterministic problems of multithreads. In `Dthread` system, it just implements threads as processes. In the executing stage, all the child processes are independent from each other. And in the commit stage, all the processes do the synchronization work. The `Dthread` library is much faster than `CoreDet` and generally as fast or even faster as `pthread`. Figure 1 describes the differences between multithreaded programs and multi-processes.

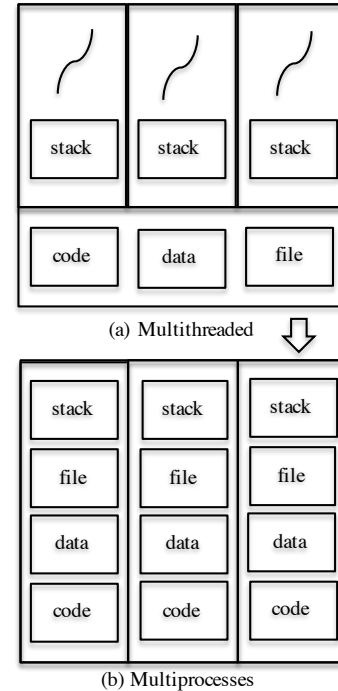


Figure 1. Change multithreaded program into multiprocess.

In this work, we also implement threads as processes, which is transparent to programmers. We provide a library,

named `libmythread`, to replace the `pthread`. Then programmers can use `-lmythread` flag when compiling their multi-threaded programs.

Outline:

The remainder of this report is organized as follows. Section 2 gives an overview of our work, including what features our library support and After that, Section ?? presents the detailed implementation, including the functions that we use when trying to support specific characteristics. Finally, Section 3 presents experimental results of `myThread`. Possible limitations and future extensions of `myThread` are discussed in Section 4.

2 Implementation

This section describes the design of `myThread`. First, we will have an overview of how `myThread` is implemented. Second, we will talk about the detail implementation from 3 aspects: How to intercept `pthread_create` to produce processes; How to share memory among the processes; How to implement synchronization.

2.1 Overview of Implementation

Figure 2 describes the overview of the whole process. The implementation has two steps. The first step is initialization, during which the library functions are wrapped and heap region, global region are also allocated in virtual address space.

To implement the goal of replacing threads with processes transparently to programmers, we should intercept all `pthread` related functions that appear in user's programs and then replace them with our own functions implemented in our library. The functions that we intercepted is shown in Table 1. The left column shows the original functions in `pthread`, and the right column shows the new functions that implemented in our `myThread` library.

As show in Figure 2, when the program calls a function in `pthread` library in run time, the funtions will be linked to the corresponding functions in our library. In this way, we can prevent `pthread_create` function from creating multiple threads, but creating multiple processes. Since multiprocess program should mimic the behavior of multithreaded programs, we should guarantee different processes sharing their heap, global, and dynamic library files. Another problem is to implement synchronization. Originally, the locks are shared among threads, but now we should initialize the locks to be shared among multiple processes. All the details of the design will be discussed in Section 2.2

2.2 Implementation Details

In this section, we will discuss the concrete technologies that we use to implement threads using processes.

Table 1. Comparasion of Library Functions

pthread functions	mythread functions
<code>pthread_create()</code>	<code>thread_create()</code>
<code>pthread_join()</code>	<code>thread_join()</code>
<code>pthread_cancel()</code>	<code>thread_cancel()</code>
<code>pthread_mutex_init()</code>	<code>thread_mutex_init()</code>
<code>pthread_mutex_lock()</code>	<code>thread_mutex_lock()</code>
<code>pthread_mutex_unlock()</code>	<code>thread_mutex_unlock()</code>
<code>pthread_mutex_trylock()</code>	<code>thread_mutex_trylock()</code>
<code>pthread_mutex_destroy()</code>	<code>thread_mutex_destroy()</code>
<code>pthread_cond_init()</code>	<code>thread_cond_init()</code>
<code>pthread_cond_wait()</code>	<code>thread_cond_wait()</code>
<code>pthread_cond_signal()</code>	<code>thread_cond_signal()</code>
<code>pthread_cond_broadcast()</code>	<code>thread_cond_broadcast()</code>
<code>pthread_cond_destroy()</code>	<code>thread_cond_destroy()</code>
<code>pthread_barrier_init()</code>	<code>thread_barrier_init()</code>
<code>pthread_barrier_wait()</code>	<code>thread_barrier_wait()</code>
<code>pthread_barrier_destroy()</code>	<code>thread_barrier_destroy()</code>

2.2.1 How to Intercept `pthread_create` to Create Processes

As mention in Section 2, our library intercepts the functions in `pthread`. So, when the programmers try to create threads, they will call the `pthread_create` function in their programs. In our implementation, we use a `clone()` system call to replace the `pthread_create()`. When we use the `clone()` function, we need to pass 4 parameters to it. The first parameter is a function pointer, which points to the function that the new process will execute. The second parameter is bottom address of a stack. We use the `mmap()` syscall to allocate a large range for the stack, caculate the end address and then pass the address to the `clone()` function. For the third parameter, we set the child process to share I/O, file system and files with the parent process. The last parameter is the arguments that the child process need, and we just use the default arguments that programmers input when they call `pthread_create` function.

To ensure that the stacks of processes can be accessed by parent process or sibling processes, we use the `MAP_SHARED` flag when we calling the `mmap()` syscall to allocate stack space. Another thing that we should mention is that we did not make the child processes to share the stack of parent process. We will leave it for future work.

2.2.2 How to Share Memory Mappings between Processes

Since we use processes to implement threads, we should make the replacement to be transparent to users. From Figure 1, we learned that all the memmory regions should be shared among the threads. That is to say, heap, global, library mappings should be shared, and there is only one copy in memory. The special case is stack. Because each thread has its own




Figure 2. Implementation overview of `myThread`. The first step is initialization when the library is loaded into memory. The second step is interception at run time.

stack, and different threads can access not only their own stacks, but also stacks of others'. So, if we try to implement multithreaded programs using processes, there should be several stacks regions in memory, and all the stack regions can be accessed by all the threads.

To implement the goals above, we tried to disable the copy-on-write characteristic of multiprocess programs. Our method is to use "cat /proc/pid/maps" command to get the memory mapping of a parent process. Before we call the clone() function, we just remapped all the memory regions to be shared except the code region. The detail process is described as Remapping Algorithm in Algorithm 1.

Algorithm 1 Remapping Algorithm

```

regions = get all the mapping regions of parent process
except code and stack
for each region ∈ regions do
    backup = mapping a new region with same SIZE;
    copy the data in region and save it in backup;
    current = remap the region area with MAP_SHARED;
    copy the data from backup to current;
end for
  
```

Since we have already talked about how to deal with stacks in Section 2. After all the above steps, we have completed most part of the tasks about implementing threads using processes. The last thing that we should do is to support synchronization, which we will talk about in detail in Section 2.2.3.

2.2.3 How to do Synchronization

In multithreaded programs, we always use lock() to define a critical region to access memory variables in an atomical way. We also use locks and conditional variables together to avoid the possible deadlock problems. So, in this section, we will discuss how to support synchronization among processes and make the multithreaded programs look like multiple processes.

We just intercept all the functions related to synchronization the same way that we deal with pthread_create. And then,

we change the attributes of locks, conditional variables, and barriers to be shared among different processes. the method that we use is shown in Table 2.

Other than lock, conditional variable and barrier, we also need to implement pthread_join to synchronize each threads. So, we just use waitpid() to wait the progresses to expire. In order to implement pthread_cancel(), we just call kill() syscall to stop a particular process.

3 Evaluation

3.1 Experimental Setup

We choose a 16-core server to do all the experiments. The server has two sockets with 256GB of memory and Intel(R) Xeon(R) CPU E5-2640 processors. For each processor, it has 256KB L1, 2MB L2, and 20MB L3 shared cache. The version of operating system is Linux3.19.0 and all applications were compiled using gcc 4.9.2 with -O2 flag.

Evaluated Applications: `myThread` is evaluated on a subset of PARSEC [2]: blacksholes, dedup, pfscan and freqmine for performance. We also get the scalability data by changing the thread numbers.

3.2 Performance

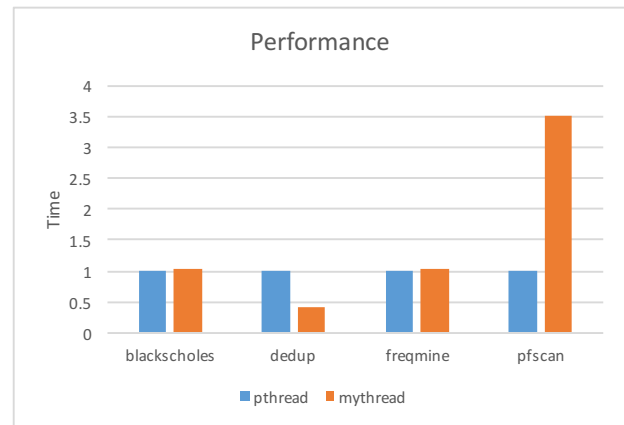


Figure 3. Performance comparison of `pthread` and `mythread` on 4 different applications.

Table 2. Synchronization Implementation

Synchronization Funtion	Actual Implementation
pthread_mutex_init(&mutex, NULL)	thread_mutexattr_init(&mutex_attr) thread_mutexattr_setpshared(&mutex_attr, PTHREAD_PROCESS_SHARED) thread_mutex_init(&mutex, &mutex_attr)
pthread_cond_init(&cond, NULL)	thread_condattr_init(&cond_attr) thread_condattr_setpshared(&cond_attr, PTHREAD_PROCESS_SHARED) thread_cond_init(&cond, &cond_attr)
pthread_barrier_init(&barrier, NULL)	thread_barrierattr_init(&barrier_attr) thread_barrierattr_setpshared(&barrier_attr, PTHREAD_PROCESS_SHARED) thread_barrier_init(&barrier, &barrier_attr)
pthread_join(tid)	waitpid((pid_t) tid)

From Figure 3, we can learn that `mythread` is comparable to `pthread` library. `mythread` performs better than `pthread` in `dedup`, but is worse in `pfscan`.

A possible reason is that isolation of memory regions among multiple threads reduces false sharing problems in `dedup` applications a lot. However, `pfscan` application may have little false sharing cases in cache.

Scalability To evaluate the scalability of `mythread`, we increase the thread number from 4 to 32, and we get the executing time of `mythread` vs. `pthread`. The performance scalability is show in Figure 4.




Figure 4. Performance scalability comparasion between `pthread` and `mythread` when increasing thread number from 4 to 32.

As show in Figure 4, we also use `pthread` as baseline. When changing the thread number from 4 to 32. The average relative executing time increase for the `mythread`. This is easy to under stand. If the thread number equals to process number, and they are both to be 1, the performance of process and thread should be the same. When we increase the thread and process number, the performance of multithreaded programs should be better at more cases. However, multiple processes can make our program more deterministic. In some precise critical applications, deterministic is even more important.

4 Limitation and Future Work

This section shows some limitations of `iReplayer` and possible extensions in the future.

Firstly, `myThread` only shares stack among child processes. Main process's stack is not shared with child processes, which means child processes can only read parent process's stack, but can not write. The semantic is not consistent with real multithreaded programs.

Secondly, `myThread` must have some potential bugs to deal with, since it can only run some of the applications in PARSEC, not all of them. To the other PARSEC applications, there are still deadlock problems.

Thirdly, comparing to `Dthread` work, we only convert the threads into processes, which is transparent to progrmers. We did not implement commit stage designed in `Dthread` paper and our evaluation work is not sufficient.

5 Acknowledgment

This work was a course project released by Dr.Tonging Liu in his operating system class, 2017. In fact, this project is a simpler version of `Dthread` project, which was conducted Dr.Liu when he pursuing his PhD in UMass Amherst. I would like to express my thanks to him for his patience and harsh training.

References

- [1] Tom Bergan, Owen Anderson, Joseph Devietti, Luis Ceze, and Dan Grossman. 2010. CoreDet: A Compiler and Runtime System for Deterministic Multithreaded Execution. In *Proceedings of the Fifteenth Edition of ASPLOS on Architectural Support for Programming Languages and Operating Systems (ASPLOS XV)*. ACM, New York, NY, USA, 53–64. <https://doi.org/10.1145/1736020.1736029>
- [2] Christian Bienia and Kai Li. 2009. PARSEC 2.0: A New Benchmark Suite for Chip-Multiprocessors. In *Proceedings of the 5th Annual Workshop on Modeling, Benchmarking and Simulation*.
- [3] Tongping Liu, Charlie Curtsinger, and Emery D. Berger. 2011. Dthreads: efficient deterministic multithreading. In *Proceedings of the Twenty-Third ACM Symposium on Operating Systems Principles (SOSP '11)*. ACM, New York, NY, USA, 327–336. <https://doi.org/10.1145/2043556.2043587>